

Project Report: Crisis Intelligence Command Center

A Multimodal RAG System for National Disaster Response

Submitted by: Aditya Sharma Date: January 22, 2026

1. Executive Summary

In the chaotic aftermath of natural disasters, decision-makers are often paralyzed not by a lack of information, but by the fragmentation of it. Critical data exists in silos: visual evidence from drones sits in one database, while urgent radio logs from ground teams sit in another. The **Crisis Intelligence Command Center** is a prototype Decision Support System (DSS) designed to bridge this cognitive gap.

Leveraging the power of **Retrieval-Augmented Generation (RAG)** and **Multimodal Vector Search**, this system creates a unified "brain" that can see, read, and remember. By encoding images and text into a shared high-dimensional vector space using **Qdrant**, the system allows responders to query complex datasets using natural language. This report details the architectural design, algorithmic logic, and societal impact of a tool built to accelerate disaster response times and save lives.

2. Problem Statement

2.1 The Societal Issue: Information Fragmentation

India is highly vulnerable to natural disasters, from the annual urban floods in Guwahati and Mumbai to cyclones in Odisha and landslides in the Himalayas. A key challenge in managing these crises is the "**fog of war**" created by disconnected data streams.

In a typical control room scenario:

- **Visual Data:** CCTV feeds, drone footage, and citizen photos are viewed on separate monitors.
- **Textual Data:** Standard Operating Procedures (SOPs), live radio transcripts, and social media distress calls are text-based.
- **Human Bottleneck:** A human operator must manually correlate a photo of a flooded street with a text report received 20 minutes earlier. This manual correlation is slow, error-prone, and impossible to scale during peak crisis moments.

2.2 Why It Matters

The inability to instantly synthesize visual and textual data leads to:

1. **Delayed Response:** Critical minutes are lost verifying if a report is true.
2. **Resource Misallocation:** Rescue boats are sent to areas that are already drained, while ignoring areas that are critically submerged.
3. **Loss of Institutional Memory:** Valuable situational awareness is often lost when a shift changes or a computer is rebooted.

Our solution addresses these issues by creating a persistent, multimodal memory that provides **instant situational awareness** to any operator, at any time.

3. System Design & Architecture

3.1 Architecture Overview

The system follows a **Dual-Stream Retrieval-Augmented Generation (RAG)** architecture. Unlike standard LLM applications that rely solely on textual training data, our system "grounds" its answers in a dynamic, real-time database of local evidence.

The architecture consists of three distinct layers:

1. **The Ingestion Layer (Sensory Input):**
 - This layer is responsible for normalizing raw data.
 - **Text Pathway:** Emergency logs and reports are processed by the [sentence-transformers](#) library.
 - **Visual Pathway:** Images (JPEG/PNG) are processed by the [CLIP](#) vision model.
 - **Storage:** Both streams are indexed into [Qdrant](#), a high-performance vector database, which acts as the system's long-term memory.
2. **The Retrieval Layer (Cognitive Search):**
 - When a user asks a question, the system does not just "guess." It performs a **semantic search**.
 - It effectively asks the database: "*Show me the 2 text reports and the 1 image most mathematically similar to this question.*"
 - This retrieval is filtered by similarity thresholds to ensure only relevant evidence is considered.
3. **The Generation Layer (Synthesis):**
 - The system uses [Google Gemini 1.5 Flash](#) as its reasoning engine.
 - The LLM receives a structured prompt containing the user's query *plus* the retrieved evidence. It then synthesizes a coherent, fact-based response.

3.2 Technical Stack & Dependencies

The solution is built on a robust, open-source Python stack designed for reproducibility and speed:

- **Frontend:** `Streamlit` (for rapid UI deployment and real-time interaction).
- **Vector Database:** `Qdrant` (Cloud/Local) for storing high-dimensional embeddings.
- **Embedding Models:**
 - Text: `all-MiniLM-L6-v2` (384 dimensions).
 - Vision: `clip-ViT-B-32` (512 dimensions).
- **LLM Provider:** `Google GenAI SDK` (Gemini 1.5 Flash).
- **Image Processing:** `Pillow (PIL)` for image manipulation.

3.3 Why Qdrant is Critical

Qdrant is the linchpin of this solution. A traditional SQL database cannot understand that "deluge" and "flood" are synonyms, nor can it match a photo of a fire to the word "blaze." Qdrant enables:

- **Vector Agnosticism:** We can store different vector sizes (384d and 512d) in parallel collections (`user_episodic_memory` and `disaster_multimodal`).
 - **Payload Filtering:** Critical for our "Safe Reset" feature, allowing us to delete specific subsets of data based on metadata tags.
 - **Speed:** Qdrant's HNSW (Hierarchical Navigable Small World) index allows for millisecond-latency searches, which is non-negotiable in disaster scenarios.
-

4. Multimodal Strategy & Vector Embeddings

4.1 The Challenge of "Cross-Modal" Search

Computers fundamentally do not understand language or images; they understand numbers. To make an AI "see" a disaster, we must translate pixels and words into lists of numbers called **Vectors**.

4.2 Text Embeddings (The Semantic Stream)

We utilize the `all-MiniLM-L6-v2` transformer model.

- **Process:** It takes a sentence like "Severe water logging at GS Road" and converts it into a dense vector of **384 floating-point numbers**.
- **Mathematical Magic:** In this 384-dimensional space, the vector for "Water logging" will be physically closer to "Flooding" than to "Earthquake." This allows the system to understand *intent* rather than just matching keywords.

4.3 Visual Embeddings (The CLIP Revolution)

We utilize OpenAI's CLIP (Contrastive Language-Image Pre-Training) model.

- **Why CLIP?** CLIP is unique because it was trained to align images and text in the *same* vector space.
 - **The Mechanism:** The vector for an *image* of a dog and the vector for the *word* "dog" will land in nearly the same spot in the 512-dimensional space.
 - **Application:** This allows our user to type "*Show me the damage*" (Text) and retrieve a photo of a landslide (Image) without any manual tagging.
-

5. Search, Memory & Recommendation Logic

5.1 The Lifecycle of a Query

The retrieval process is a sophisticated, multi-step algorithm designed to minimize noise and maximize relevance:

1. **User Input:** The commander types: "*Is there flooding in sector 4?*"
2. **Dual-Encoding:**
 - The query is encoded into a **384d vector** (for text search).
 - The query is encoded into a **512d vector** (for image search).
3. **Parallel Execution:** Both vectors are sent to Qdrant simultaneously.
4. **Threshold Filtering:**
 - **Text Filter (>0.40):** We ignore any text logs with a similarity score below 0.40. This prevents the AI from hallucinating connections where none exist.
 - **Visual Filter (>0.25):** We retrieve the top image only if it meets a confidence threshold of 0.25.
 - **Negative Constraints:** A logic layer checks for keywords like "Don't show me..." to suppress visual output if requested.
5. **Synthesis:** The valid results are stitched into a "Context Block" and sent to Gemini 1.5 Flash for the final answer.

5.2 Memory Management: The "Safe Reset" Protocol

A major innovation in this system is its lifecycle management of memory. In a continuous operation, a system can become cluttered with old data.

- **Episodic Memory (Short-Term):** Every user question and AI answer is stored with the metadata tag `role="user"` or `role="assistant"`.
- **Semantic Memory (Long-Term):** The core disaster reports are stored with the tag `role="system_report"`.

The Reset Mechanism: When the user clicks "**Start New Scenario**", the system executes a targeted delete command:

```
DELETE FROM user_episodic_memory WHERE role IN ['user', 'assistant']
```

- This surgically wipes the conversation history while leaving the critical "System Reports" intact, allowing the operator to start fresh without reloading the database.
-

Based on the specific dataset you have ingested (Pan-India disaster logs and the matched images), here is the **Simulated Operational Scenario** section for your report.

You can copy-paste this directly into **Section 6** of your PDF report. It creates a coherent narrative that proves your system correctly links specific text logs to their corresponding visual evidence.

6. Simulated Operational Scenario

To demonstrate the system's "Grounding" capabilities, we present a walkthrough of four distinct crisis events detected by the Command Center using the ingested dataset.

Phase 1: Structural Integrity Assessment (Gujarat Sector)

- **Commander Query:** *"Report on the structural collapse in Bhuj."*
- **System Action (Text Retrieval):** The system queries the `user_episodic_memory` collection and retrieves the log:
"Gujarat Control Room: Structural collapse in Bhuj old town; masonry collapse of heritage structure, collapse zone cordoned..."
- **System Action (Visual Retrieval):** Simultaneously, the system queries `disaster_multimodal` and identifies a matching image file:
`bhuj-gujarat-india-january-a-ruined-stone-structure...jpg`.
- **AI Synthesis:**
"Confirmed structural failure in Bhuj Old Town. The retrieved visual evidence shows a multi-story masonry structure with significant facade collapse. Rubble patterns suggest a total loss of the front elevation. Urban Search and Rescue (USAR) canine teams are advised for void searching."

Phase 2: Critical Infrastructure Failure (Karnataka Sector)

- **Commander Query:** *"Show me the road damage report for Yellapur."*
- **System Action (Text Retrieval):**
"Karnataka SDRF: Flash flood in Yellapur taluk; bridge SC-B collapsed at span 3, two vehicles submerged..."
- **System Action (Visual Retrieval):** The system matches the location "Yellapur" to the file
`dangerous-landslide-in-road-of-yellapur-karnataka-india.jpg`.
- **AI Synthesis:**
"Major infrastructure severance detected. Visual evidence confirms a complete washout of the roadway due to a landslide/flash flood event. The red soil

substructure is exposed, indicating deep slope failure. The route is impassable for relief trucks; alternate routes via Hubli are recommended."

Phase 3: Mass Casualty Incident (Rail Transport)

- **Commander Query:** "*What is the situation with the train derailment?*"
- **System Action (Text Retrieval):**
"Bihar NDRF Liaison: Train derailment near Hajipur due to track washout; multiple coaches overturned, extrication priority for trapped passengers..."
- **System Action (Visual Retrieval):** The system identifies the semantic match for "train wreck" in the file
[`-kanpur-uttar-pradesh-india-november-rescue-and-relief-works.. .jpg.`](#)
- **AI Synthesis:**
"Severe rail incident reported. Text logs indicate track washout as the cause. Visual evidence shows multiple blue railway coaches capsized off the embankment with heavy crowding nearby. Heavy hydraulic cutters and cranes are required immediately for extrication."

Phase 4: Urban Flood Response (Assam/Kerala Context)

- **Commander Query:** "*Status of rescue operations in the flood zones?*"
- **System Action (Text Retrieval):**
"Kerala Control Room: Emergency water-rescue ongoing in Idukki... 4 boats deployed..."
- **System Action (Visual Retrieval):** The system retrieves
[`guwahati-india-october-national-disaster-response-force-ndrf-p ersonnel-rescue-people-.jpg.`](#)
- **AI Synthesis:**
"Active rescue operations are underway. Visuals show NDRF personnel utilizing orange inflatable boats to navigate urban streets where water levels have reached waist-height. Current priority is evacuating civilians from inundated ground-floor residences."

7. Limitations & Ethics

7.1 Known Failure Modes

- **Semantic Ambiguity:** In scenarios with sparse data, the model might "force" a match. For example, if asked for "fire" but the database only has "flood" images, a low threshold might incorrectly retrieve a flood image. Strict threshold tuning is required.
- **Latency Dependency:** The system relies on external API calls to Google Gemini. In a grid-down scenario where internet connectivity is spotty, the inference engine would fail, though the local vector database might still be accessible.

7.2 Ethical Considerations

- **Algorithmic Bias:** The CLIP model is trained on internet-scale data, which often underrepresents disasters in the Global South. This could lead to lower retrieval accuracy for housing types or landscapes specific to rural India compared to Western cities.
 - **Data Privacy:** Uploading images of disaster victims to cloud-based AI services poses a privacy risk. A production version of this tool must comply with the **Digital Personal Data Protection (DPDP) Act, 2023**, ensuring that sensitive data is anonymized or processed on local servers.
 - **Human-in-the-Loop:** This AI is a tool for *assistance*, not *automation*. It should never be given the authority to deploy resources autonomously. All AI suggestions must be verified by a human commander.
-

8. Conclusion & Future Scope

The **Crisis Intelligence Command Center** demonstrates that modern AI can be more than just a chatbot—it can be a reliable partner in high-stakes environments. By successfully integrating Multimodal RAG with robust state management, we have created a prototype that solves the critical problem of information fragmentation.

Future Roadmap:

1. **GIS Integration:** Plotting retrieved data points on a live map interface (e.g., Leaflet or Google Maps).
2. **Voice Interface:** allowing commanders to query the system via radio voice commands (Speech-to-Text).
3. **Local Deployment:** Migrating from Gemini API to a locally hosted Llama 3 model to ensure operation during internet blackouts.

This project serves as a foundational step toward a smarter, faster, and more resilient national disaster response infrastructure.